

Experiment No. 2
Analyze the Titanic Survival Dataset and apply appropriate regression technique
Date of Performance: 31-7-23
Date of Submission: 7-8-23

Aim: Analyze the Titanic Survival Dataset and apply appropriate Regression Technique.

Objective: Able to perform various feature engineering tasks, apply logistic regression on the given dataset and maximize the accuracy.

Theory:


Logistic Regression was used in the biological sciences in early twentieth century. It was then used in many social science applications. Logistic Regression is used when the dependent variable(target) is categorical and is binary in nature. In order to perform binary classification the logistic regression techniques makes use of Sigmoid function.

For example,

To predict whether an email is spam (1) or (0)

Whether the tumor is malignant (1) or not (0)

Consider a scenario where we need to classify whether an email is spam or not. If we use linear regression for this problem, there is a need for setting up a threshold based on which classification can be done. Say if the actual class is malignant, predicted continuous value 0.4 and the threshold value is 0.5, the data point will be classified as not malignant which can lead to serious consequence in real time.



From this example, it can be inferred that linear regression is not suitable for classification problem. Linear regression is unbounded, and this brings logistic regression into picture. Their value strictly ranges from 0 to 1.

Dataset:

The sinking of the Titanic is one of the most infamous shipwrecks in history.

On April 15, 1912, during her maiden voyage, the widely considered “unsinkable” RMS Titanic sank after colliding with an iceberg. Unfortunately, there weren’t enough lifeboats for everyone onboard, resulting in the death of 1502 out of 2224 passengers and crew.

While there was some element of luck involved in surviving, it seems some groups of people were more likely to survive than others.

In this challenge, we ask you to build a predictive model that answers the question: “what sorts of people were more likely to survive?” using passenger data (ie name, age, gender, socioeconomic class, etc).

Variable	Definition	Key
survival	Survival	0 = No, 1 = Yes
pclass	Ticket class	1 = 1st, 2 = 2nd, 3 = 3rd
sex	Sex	
Age	Age in years	
sibsp	# of siblings / spouses aboard the Titanic	
parch	# of parents / children aboard the Titanic	
ticket	Ticket number	
fare	Passenger fare	
cabin	Cabin number	

embarked	Port of Embarkation	C = Cherbourg, Q = Queenstown, S = Southampton
----------	---------------------	--

Variable Notes pclass: A proxy for socio-economic status (SES)

1st = Upper, 2nd = Middle, 3rd = Lower

age: Age is fractional if less than 1. If the age is estimated, is it in the form of xx.5

sibsp: The dataset defines family relations in this way...

Sibling = brother, sister, stepbrother, stepsister

Spouse = husband, wife (mistresses and fiancés were ignored)

parch: The dataset defines family relations in this way...

Parent = mother, father

Child = daughter, son, stepdaughter, stepson

Some children travelled only with a nanny, therefore parch=0 for them.

Code:

Conclusion:

The following features have been chosen for the model development process: passenger class (Pclass), gender (Sex), age (Age), number of parents and children (Parch), and number of siblings and spouses aboard (SibSp). These characteristics are critical because they may affect survival rates and reflect socioeconomic factors, including the precedence accorded to

passengers in higher class, women during evacuations, children and the elderly receiving priority, the presence of family members, and associations between the departure port and socioeconomic backgrounds.

With an accuracy score of 0.8076 for the training set, the model can accurately predict results of survival for this dataset. The accuracy of the test results is 0.7821, indicating that the model functions well with fresh, untested data. These scores for accuracy indicate that the model successfully generalises to unfamiliar data

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
import warnings
```

```
data = pd.read_csv("/content/train (2).csv")
```

```
print(data)
```

```

   PassengerId  Survived  Pclass \
0             1         0       3
1             2         1       1
2             3         1       3
3             4         1       1
4             5         0       3
..          ...         ...     ...
886          887         0       2
887          888         1       1
888          889         0       3
889          890         1       1
890          891         0       3
```

```

                                Name    Sex  Age  SibSp \
0      Braund, Mr. Owen Harris    male  22.0      1
1  Cumings, Mrs. John Bradley (Florence Briggs Th... female  38.0      1
2      Heikkinen, Miss. Laina    female  26.0      0
3  Futrelle, Mrs. Jacques Heath (Lily May Peel)    female  35.0      1
4      Allen, Mr. William Henry    male  35.0      0
..          ...         ...     ...
886      Montvila, Rev. Juozas    male  27.0      0
887      Graham, Miss. Margaret Edith    female  19.0      0
888  Johnston, Miss. Catherine Helen "Carrie"    female   NaN      1
889      Behr, Mr. Karl Howell    male  26.0      0
890      Dooley, Mr. Patrick    male  32.0      0
```

```

   Parch  Ticket  Fare Cabin Embarked
0      0   A/5 21171   7.2500   NaN      S
1      0   PC 17599  71.2833   C85      C
2      0  STON/O2. 3101282  7.9250   NaN      S
3      0   113803  53.1000  C123      S
4      0   373450  8.0500   NaN      S
..    ...         ...     ...     ...
886      0   211536  13.0000   NaN      S
887      0   112053  30.0000  B42      S
888      2   W./C. 6607  23.4500   NaN      S
889      0   111369  30.0000  C148      C
890      0   370376  7.7500   NaN      Q
```

```
[891 rows x 12 columns]
```

```
data.shape
```

```
(891, 12)
```

```
data.info() # getting some informations about the data
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column      Non-Null Count  Dtype
---  -
0   PassengerId  891 non-null    int64
1   Survived     891 non-null    int64
2   Pclass       891 non-null    int64
3   Name         891 non-null    object
4   Sex          891 non-null    object
5   Age          714 non-null    float64
6   SibSp        891 non-null    int64
7   Parch        891 non-null    int64
8   Ticket       891 non-null    object
9   Fare         891 non-null    float64
10  Cabin        204 non-null    object
11  Embarked     889 non-null    object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

```
data.isnull().sum() # check the number of missing values in each column
```

```

PassengerId    0
Survived        0
Pclass         0
Name           0
Sex            0
Age           177
SibSp          0
Parch          0
Ticket         0
Fare           0
Cabin         687
Embarked       2
dtype: int64

```

```
data = data.drop(columns='Cabin', axis=1)
```

```
data['Age'].fillna(data['Age'].mean(), inplace=True) # replacing the missing values in "Age" column with mean value
```

```
print(data['Embarked'].mode()) # finding the mode value of "Embarked" column
```

```

0    S
Name: Embarked, dtype: object

```

```
print(data['Embarked'].mode()[0])
```

```
S
```

```
data['Embarked'].fillna(data['Embarked'].mode()[0], inplace=True) # replacing the missing values in "Embarked" column with mode value
```

```
data.isnull().sum() # check the number of missing values in each column
```

```

PassengerId    0
Survived        0
Pclass         0
Name           0
Sex            0
Age            0
SibSp          0
Parch          0
Ticket         0
Fare           0
Embarked       0
dtype: int64

```

```
data.describe()
```

	PassengerId	Survived	Pclass	Age	SibSp	Parch	Fare
count	891.000000	891.000000	891.000000	891.000000	891.000000	891.000000	891.000000
mean	446.000000	0.383838	2.308642	29.699118	0.523008	0.381594	32.204208
std	257.353842	0.486592	0.836071	13.002015	1.102743	0.806057	49.693429
min	1.000000	0.000000	1.000000	0.420000	0.000000	0.000000	0.000000
25%	223.500000	0.000000	2.000000	22.000000	0.000000	0.000000	7.910400
50%	446.000000	0.000000	3.000000	29.699118	0.000000	0.000000	14.454200
75%	668.500000	1.000000	3.000000	35.000000	1.000000	0.000000	31.000000
max	891.000000	1.000000	3.000000	80.000000	8.000000	6.000000	512.329200

```
data['Survived'].value_counts() # finding the number of people survived and not survived
```

```

0    549
1    342
Name: Survived, dtype: int64

```

```
data['Sex'].value_counts()
```

```

male    577
female  314
Name: Sex, dtype: int64

```


```
# number of survivors Gender wise
```

```
# 1st male and other female
```

```
# 0 are the one who did not survived
```

```
sns.countplot(x='Sex', hue='Survived', data=data)
```

<Axes: xlabel='Sex', ylabel='count'>



```
data['Embarked'].value_counts()
```

```
S    646
C    168
Q     77
Name: Embarked, dtype: int64
```

```
# converting categorical Columns
```

```
data.replace({'Sex':{'male':0,'female':1}, 'Embarked':{'S':0,'C':1,'Q':2}}, inplace=True)
```

```
X = data.drop(columns = ['PassengerId','Name','Ticket','Survived'],axis=1)
```

```
Y = data['Survived']
```

```
print(X)
```

	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
0	3	0	22.000000	1	0	7.2500	0
1	1	1	38.000000	1	0	71.2833	1
2	3	1	26.000000	0	0	7.9250	0
3	1	1	35.000000	1	0	53.1000	0
4	3	0	35.000000	0	0	8.0500	0
..
886	2	0	27.000000	0	0	13.0000	0
887	1	1	19.000000	0	0	30.0000	0
888	3	1	29.699118	1	2	23.4500	0
889	1	0	26.000000	0	0	30.0000	1
890	3	0	32.000000	0	0	7.7500	2

```
[891 rows x 7 columns]
```

```
print(Y)
```

```
0    0
1    1
2    1
3    1
4    0
..
886  0
887  1
888  0
889  1
890  0
Name: Survived, Length: 891, dtype: int64
```

```
data.head()
```



```

PassengerId  Survived  Pclass
X.head()

Pclass  Sex  Age  SibSp  Parch  Fare  Embarked
0      3    0  22.0    1     0   7.2500      0
1      1    1  38.0    1     0  71.2833      1
2      3    1  26.0    0     0   7.9250      0
3      1    1  35.0    1     0  53.1000      0
4      3    0  35.0    0     0   8.0500      0

```

```
Y.head()
```

```

0    0
1    1
2    1
3    1
4    0

```

```
Name: Survived, dtype: int64
```

```
#Splitting the data into training data & Test data
```

```
X_train, X_test, Y_train, Y_test = train_test_split(X,Y, test_size=0.2, random_state=2)
```

```
print(X.shape, X_train.shape, X_test.shape)
```

```
(891, 7) (712, 7) (179, 7)
```

```
logr = LogisticRegression()
```

```
# training the Logistic Regression model with training data
```

```
logr.fit(X_train, Y_train)
```

```

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: Conver
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

```
Increase the number of iterations (max_iter) or scale the data as shown in:
```

```
https://scikit-learn.org/stable/modules/preprocessing.html
```

```
Please also refer to the documentation for alternative solver options:
```

```
https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regression
```

```
n_iter_i = _check_optimize_result(
```

```
    LogisticRegression
```

```
LogisticRegression())
```

```
# accuracy on training data
```

```
X_train_prediction = logr.predict(X_train)
```

```
training_data_accuracy = accuracy_score(Y_train, X_train_prediction)
```

```
print("Accuracy score of training data:", training_data_accuracy)
```

```
Accuracy score of training data: 0.8075842696629213
```

```
# accuracy on test data
```

```
X_test_prediction = logr.predict(X_test)
```

```
test_data_accuracy = accuracy_score(Y_test, X_test_prediction)
```

```
print('Accuracy score of test data : ', test_data_accuracy)
```

```
Accuracy score of test data : 0.7821229050279329
```