



Experiment No. 4
Apply Random Forest Algorithm on Adult Census Income Dataset and analyze the performance of the model
Date of Performance: 14-8-23
Date of Submission: 22-8-23



Aim: Apply Random Forest Algorithm on Adult Census Income Dataset and analyze the performance of the model.

Objective: Able to perform various feature engineering tasks, apply Random Forest Algorithm on the given dataset and maximize the accuracy, Precision, Recall, F1 score.

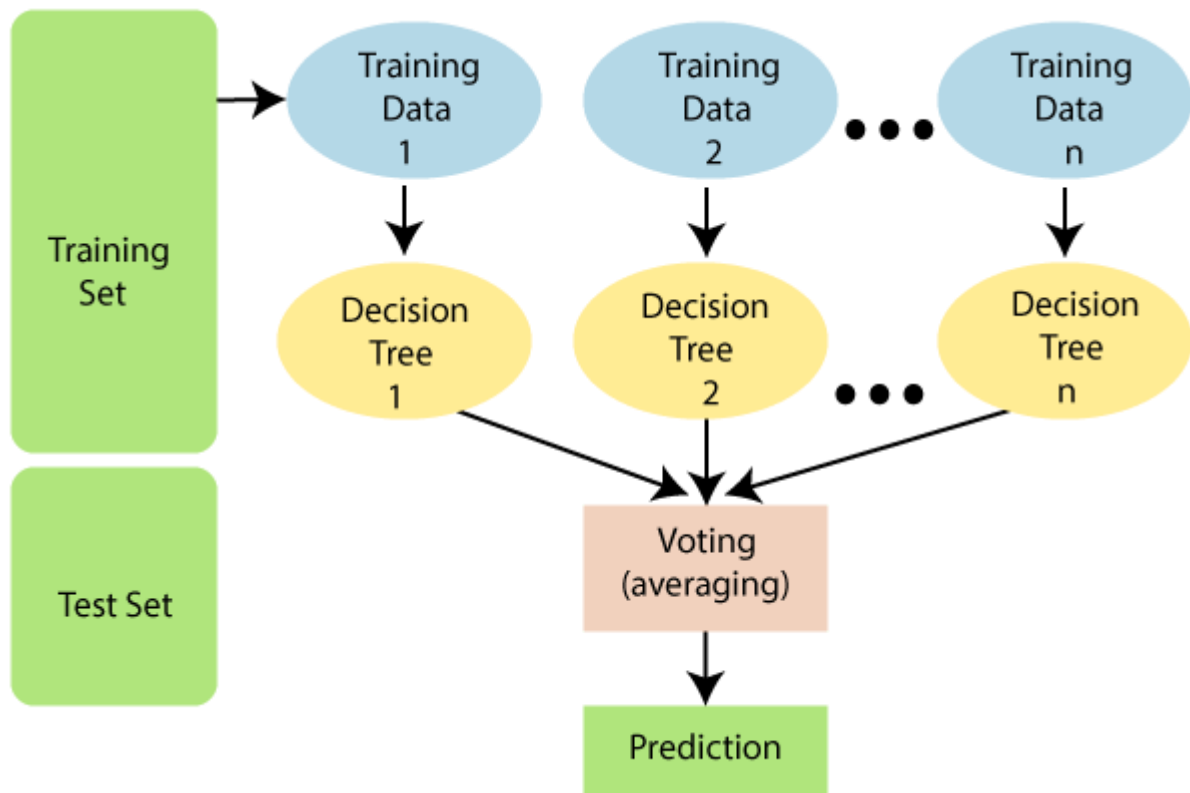
Theory:

Random Forest is a popular machine learning algorithm that belongs to the supervised learning technique. It can be used for both Classification and Regression problems in ML. It is based on the concept of ensemble learning, which is a process of combining multiple classifiers to solve a complex problem and to improve the performance of the model.

As the name suggests, "Random Forest is a classifier that contains a number of decision trees on various subsets of the given dataset and takes the average to improve the predictive accuracy of that dataset." Instead of relying on one decision tree, the random forest takes the prediction from each tree and based on the majority votes of predictions, and it predicts the final output.

The greater number of trees in the forest leads to higher accuracy and prevents the problem of overfitting.

The below diagram explains the working of the Random Forest algorithm:



Dataset:

Predict whether income exceeds \$50K/yr based on census data. Also known as "Adult" dataset.

Attribute Information:

Listing of attributes:

>50K, <=50K. age:

continuous.

workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov,

Without-pay, Never-worked. fnlwgt: continuous.



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.

education-num: continuous.

marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Marriedspouse-absent, Married-AF-spouse.

occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.

relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.

race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black. sex: Female,

Male. capital-gain: continuous. capital-loss: continuous. hours-per-week:

continuous.

native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, OutlyingUS(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad & Tobago, Peru, Hong, Holand-Netherlands.

Code:

Conclusion:

These variables don't have many significant linear relationships with one another; instead, their correlations are typically weaker. Age shows a marginally favourable link with the quantity of schooling and weekly hours worked. There is a little positive association between education levels and capital gains. Additionally, there is a somewhat negative association between capital losses and profits.



Accuracy: 85.44% of the time, the model accurately predicts income levels in the majority of cases.

True positives (8015), false positives (628), and false negatives (1047) are shown in the confusion matrix. Actual negatives (1819) forecasts.

Precision: 0.74 for income number one and 0.08 for income number zero.

Recall: 0.93 is the recall for income 0 and 0.63 is the recall for income 1.

F1-score: A measure of overall model efficacy, the F1-score is the mean of precision and recall. There are 0 for 0.91 and 1 for 0.68 in it.

Generally speaking, Random Forest produces superior outcomes to Decision Trees. Multiple Decision Tree forecasts are combined in the Random Forest model, which can increase generalisation and accuracy.

```
# Import libraries
import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
sns.set(style='white', context='notebook', palette='deep')
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV, cross_val_score, StratifiedKFold, learning_curve, train_test_split, KFold
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score

# To ignore warning messages
import warnings
warnings.filterwarnings('ignore')
# Adult dataset path
adult_dataset_path = "/content/adult.csv"
# Function for loading adult dataset
def load_adult_data(adult_path=adult_dataset_path):
    csv_path = os.path.join(adult_path)
    return pd.read_csv(csv_path)

# Calling load adult function and assigning to a new variable df
df = load_adult_data()
# load top 3 rows values from adult dataset
df.head(3)
```

	age	workclass	fnlwgt	education	educational-num	marital-status	occupation	relationship	race	gender	capital-gain
0	25	Private	226802	11th	7	Never-married	Machine-op-inspct	Own-child	Black	Male	0
1	38	Private	89814	HS-grad	9	Married-civ-spouse	Farming-fishing	Husband	White	Male	0
2	28	Local-gov	336951	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Husband	White	Male	0

```
print ("Rows : " ,df.shape[0])
print ("Columns : " ,df.shape[1])
print ("\nFeatures : \n" ,df.columns.tolist())
print ("\nMissing values : ", df.isnull().sum().values.sum())
print ("\nUnique values : \n",df.nunique())

Rows : 48842
Columns : 15

Features :
['age', 'workclass', 'fnlwgt', 'education', 'educational-num', 'marital-status', 'occupation', 'relationship', 'race', 'gender', 'capital-gain', 'capital-loss', 'hours-per-week', 'native-country', 'income']

Missing values : 0

Unique values :
age          74
workclass     9
fnlwgt      28523
education    16
educational-num 16
marital-status 7
occupation   15
relationship  6
race         5
gender       2
capital-gain 123
capital-loss  99
hours-per-week 96
native-country 42
income       2
dtype: int64
```

```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 48842 entries, 0 to 48841
Data columns (total 15 columns):
#   Column              Non-Null Count  Dtype
---  -
0   age                 48842 non-null  int64
1   workclass           48842 non-null  object
2   fnlwgt              48842 non-null  int64
3   education           48842 non-null  object
```

```
4 educational-num 48842 non-null int64
5 marital-status 48842 non-null object
6 occupation      48842 non-null object
7 relationship    48842 non-null object
8 race            48842 non-null object
9 gender          48842 non-null object
10 capital-gain    48842 non-null int64
11 capital-loss   48842 non-null int64
12 hours-per-week 48842 non-null int64
13 native-country 48842 non-null object
14 income         48842 non-null object
dtypes: int64(6), object(9)
memory usage: 5.6+ MB
```

```
df.describe()
```

	age	fnlwgt	educational-num	capital-gain	capital-loss	hours-per-week
count	48842.000000	4.884200e+04	48842.000000	48842.000000	48842.000000	48842.000000
mean	38.643585	1.896641e+05	10.078089	1079.067626	87.502314	40.422382
std	13.710510	1.056040e+05	2.570973	7452.019058	403.004552	12.391444
min	17.000000	1.228500e+04	1.000000	0.000000	0.000000	1.000000
25%	28.000000	1.175505e+05	9.000000	0.000000	0.000000	40.000000
50%	37.000000	1.781445e+05	10.000000	0.000000	0.000000	40.000000
75%	48.000000	2.376420e+05	12.000000	0.000000	0.000000	45.000000
max	90.000000	1.490400e+06	16.000000	99999.000000	4356.000000	99.000000

```
df.head()
```

	age	workclass	fnlwgt	education	educational-num	marital-status	occupation	relationship	race	gender	capital-gain
0	25	Private	226802	11th	7	Never-married	Machine-op-inspct	Own-child	Black	Male	0
1	38	Private	89814	HS-grad	9	Married-civ-spouse	Farming-fishing	Husband	White	Male	0
2	28	Local-gov	336951	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Husband	White	Male	0
3	44	Private	160323	Some-college	10	Married-civ-spouse	Machine-op-inspct	Husband	Black	Male	7688
4	18	?	103497	Some-college	10	Never-married	?	Own-child	White	Female	0

```
# checking "?" total values present in particular 'workclass' feature
df_check_missing_workclass = (df['workclass']=='?').sum()
df_check_missing_workclass

2799
```

```
# checking "?" total values present in particular 'occupation' feature
df_check_missing_occupation = (df['occupation']=='?').sum()
df_check_missing_occupation

2809
```

```
# checking "?" values, how many are there in the whole dataset
df_missing = (df=='?').sum()
df_missing
```

age	0
workclass	2799
fnlwgt	0
education	0
educational-num	0
marital-status	0
occupation	2809
relationship	0
race	0
gender	0
capital-gain	0
capital-loss	0
hours-per-week	0
native-country	857
income	0
dtype:	int64

```
percent_missing = (df=='?').sum() * 100/len(df)
percent_missing
```

```

age          0.000000
workclass    5.730724
fnlwgt       0.000000
education    0.000000
educational-num 0.000000
marital-status 0.000000
occupation   5.751198
relationship 0.000000
race         0.000000
gender       0.000000
capital-gain 0.000000
capital-loss 0.000000
hours-per-week 0.000000
native-country 1.754637
income       0.000000
dtype: float64

```

```

# find total number of rows which doesn't contain any missing value as '?'
df.apply(lambda x: x != '?',axis=1).sum()

```

```

age          48842
workclass    46043
fnlwgt       48842
education    48842
educational-num 48842
marital-status 48842
occupation   46033
relationship 48842
race         48842
gender       48842
capital-gain 48842
capital-loss 48842
hours-per-week 48842
native-country 47985
income       48842
dtype: int64

```

```

# dropping the rows having missing values in workclass
df = df[df['workclass'] != '?']
df.head()

```

	age	workclass	fnlwgt	education	educational-num	marital-status	occupation	relationship	race	gender	capital-gain
0	25	Private	226802	11th	7	Never-married	Machine-op-inspct	Own-child	Black	Male	0
1	38	Private	89814	HS-grad	9	Married-civ-spouse	Farming-fishing	Husband	White	Male	0
2	28	Local-gov	336951	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Husband	White	Male	0
3	44	Private	160323	Some-college	10	Married-civ-spouse	Machine-op-inspct	Husband	Black	Male	7688
5	34	Private	198693	10th	6	Never-married	Other-service	Not-in-family	White	Male	0

```

# select all categorical variables
df_categorical = df.select_dtypes(include=['object'])
# checking whether any other column contains '?' value
df_categorical.apply(lambda x: x=='?',axis=1).sum()

```

```

workclass    0
education    0
marital-status 0
occupation   10
relationship  0
race         0
gender       0
native-country 811
income       0
dtype: int64

```

```

from sklearn import preprocessing
# encode categorical variables using label Encoder
# select all categorical variables
df_categorical = df.select_dtypes(include=['object'])
df_categorical.head()

```



```
workclass    education    marital-status    occupation    relationship    race    gender    native-country    income
0  Private      10th      Married      Private      Self-employed-not-inc.    White      Male      United-States    2156
# apply label encoder to df_categorical
le = preprocessing.LabelEncoder()
df_categorical = df_categorical.apply(le.fit_transform)
df_categorical.head()
```

	workclass	education	marital-status	occupation	relationship	race	gender	native-country	income
0	2	1	4	6	3	2	1	39	0
1	2	11	2	4	0	4	1	39	0
2	1	7	2	10	0	4	1	39	1
3	2	15	2	6	0	2	1	39	1
5	2	0	4	7	1	4	1	39	0

```
# Next, Concatenate df_categorical dataframe with original df (dataframe)
# first, Drop earlier duplicate columns which had categorical values
df = df.drop(df_categorical.columns,axis=1)
df = pd.concat([df,df_categorical],axis=1)
df.head()
```


	age	fnlwgt	educational-num	capital-gain	capital-loss	hours-per-week	workclass	education	marital-status	occupation	relationship
0	25	226802	7	0	0	40	2	1	4	6	
1	38	89814	9	0	0	50	2	11	2	4	
2	28	336951	12	0	0	40	1	7	2	10	
3	44	160323	10	7688	0	40	2	15	2	6	
5	34	198693	6	0	0	30	2	0	4	7	

```
# look at column type
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 46033 entries, 0 to 48841
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   46033 non-null  int64
1   fnlwgt                46033 non-null  int64
2   educational-num       46033 non-null  int64
3   capital-gain          46033 non-null  int64
4   capital-loss          46033 non-null  int64
5   hours-per-week        46033 non-null  int64
6   workclass             46033 non-null  int64
7   education             46033 non-null  int64
8   marital-status        46033 non-null  int64
9   occupation            46033 non-null  int64
10  relationship          46033 non-null  int64
11  race                  46033 non-null  int64
12  gender                46033 non-null  int64
13  native-country        46033 non-null  int64
14  income                46033 non-null  int64
dtypes: int64(15)
memory usage: 5.6 MB
```

```
plt.figure(figsize=(14,10))
sns.heatmap(df.corr(),annot=True,fmt='.2f')
plt.show()
```





```
# convert target variable income to categorical
df['income'] = df['income'].astype('category')
# check df info again whether everything is in right format or not
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 46033 entries, 0 to 48841
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                    46033 non-null  int64
1   fnlwgt                 46033 non-null  int64
2   educational-num        46033 non-null  int64
3   capital-gain           46033 non-null  int64
4   capital-loss           46033 non-null  int64
5   hours-per-week         46033 non-null  int64
6   workclass              46033 non-null  int64
7   education              46033 non-null  int64
8   marital-status         46033 non-null  int64
9   occupation             46033 non-null  int64
10  relationship            46033 non-null  int64
11  race                   46033 non-null  int64
12  gender                 46033 non-null  int64
13  native-country         46033 non-null  int64
14  income                 46033 non-null  category
dtypes: category(1), int64(14)
memory usage: 5.3 MB
```

```
# Importing train_test_split
from sklearn.model_selection import train_test_split
# Putting independent variables/features to X
X = df.drop('income',axis=1)
# Putting response/dependent variable/feature to y
y = df['income']
X.head(3)
```

	age	fnlwgt	educational- num	capital- gain	capital- loss	hours- per- week	workclass	education	marit sta
0	25	226802	7	0	0	40	2	1	
1	38	89814	9	0	0	50	2	11	

```
y.head(3)
```

```
0    0
1    0
2    1
```

```
Name: income, dtype: category
Categories (2, int64): [0, 1]

# Splitting the data into train and test
X_train,X_test,y_train,y_test = train_test_split(X,y)
X_train.head()
```

	age	fnlwgt	educational-num	capital-gain	capital-loss	hours-per-week	workclass	education	m
13554	58	196502	10	0	0	60	2	15	
46282	27	297457	9	0	0	40	2	11	
25679	27	30244	9	0	0	80	4	11	
8775	42	165309	9	0	0	50	2	11	

```
test_size = 0.20
seed = 7
num_folds = 10
scoring = 'accuracy'
# Params for Random Forest
num_trees = 100
max_features = 3

random_forest = RandomForestClassifier(n_estimators=250,max_features=5)
random_forest.fit(X_train, y_train)
predictions = random_forest.predict(X_test)
print("Accuracy: %s%" % (100*accuracy_score(y_test, predictions)))
print(confusion_matrix(y_test, predictions))
print(classification_report(y_test, predictions))
```

```
Accuracy: 85.44617256060475%
[[8015 628]
 [1047 1819]]
```

	precision	recall	f1-score	support
0	0.88	0.93	0.91	8643
1	0.74	0.63	0.68	2866
accuracy			0.85	11509
macro avg	0.81	0.78	0.80	11509
weighted avg	0.85	0.85	0.85	11509